

# import library

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
import math
import string
import nltk
from nltk.stem import PorterStemmer
import warnings #But we need to hide these warnings
warnings.filterwarnings('ignore')
warnings.filterwarnings('ignore', category=DeprecationWarning)
warnings.filterwarnings("ignore",category=UserWarning)
sns.set_style("whitegrid")
%matplotlib inline
np.random.seed(7)
```

```
In [2]: df=pd.read_csv('C:/Users/User/Desktop/PROJECTS/Amazon Product Review/1429_1.csv')
print(df)
```

	id	\
0	AVqkIhwDv8e3D10-lebb	
1	AVqkIhwDv8e3D10-lebb	
2	AVqkIhwDv8e3D10-lebb	
3	AVqkIhwDv8e3D10-lebb	
4	AVqkIhwDv8e3D10-lebb	
...	...	...
34655	AVpfiBlyLJeJML43-4Tp	
34656	AVpfiBlyLJeJML43-4Tp	
34657	AVpfiBlyLJeJML43-4Tp	
34658	AVpfiBlyLJeJML43-4Tp	
34659	AVpfiBlyLJeJML43-4Tp	

	name	asins	brand	\
0	All-New Fire HD 8 Tablet, 8 HD Display, Wi-Fi,...	B01AHB9CN2	Amazon	
1	All-New Fire HD 8 Tablet, 8 HD Display, Wi-Fi,...	B01AHB9CN2	Amazon	
2	All-New Fire HD 8 Tablet, 8 HD Display, Wi-Fi,...	B01AHB9CN2	Amazon	
3	All-New Fire HD 8 Tablet, 8 HD Display, Wi-Fi,...	B01AHB9CN2	Amazon	
4	All-New Fire HD 8 Tablet, 8 HD Display, Wi-Fi,...	B01AHB9CN2	Amazon	

```
In [3]: print(len(df))
```

34660

```
In [4]: df.shape
```

Out[4]: (34660, 21)

```
In [5]: df.columns
```

Out[5]: Index(['id', 'name', 'asins', 'brand', 'categories', 'keys', 'manufacturer', 'reviews.date', 'reviews.dateAdded', 'reviews.dateSeen', 'reviews.didPurchase', 'reviews.doRecommend', 'reviews.id', 'reviews.numHelpful', 'reviews.rating', 'reviews.sourceURLs', 'reviews.text', 'reviews.title', 'reviews.userCity', 'reviews.userProvince', 'reviews.username'], dtype='object')

```
In [6]: df.head()
```

Out[6]:

	id	name	asins	brand	categories	keys	manufac
0	AVqkIhwDv8e3D10-lebb	All-New Fire HD 8 Tablet, 8 HD Display, Wi-Fi,...	B01AHB9CN2	Amazon	Electronics,iPad & Tablets,All Tablets,Fire Ta...	841667104676,amazon/53004484,amazon/b01ahb9cn2...	An
1	AVqkIhwDv8e3D10-lebb	All-New Fire HD 8 Tablet, 8 HD Display, Wi-Fi,...	B01AHB9CN2	Amazon	Electronics,iPad & Tablets,All Tablets,Fire Ta...	841667104676,amazon/53004484,amazon/b01ahb9cn2...	An
	AVqkIhwDv8e3D10-lebb	All-New Fire HD 8 Tablet, 8 HD Display, Wi-Fi,...	B01AHB9CN2	Amazon	Electronics,iPad & Tablets,All Tablets,Fire Ta...	841667104676,amazon/53004484,amazon/b01ahb9cn2...	An

In [7]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 34660 entries, 0 to 34659
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                    34660 non-null  object
1   name                                27900 non-null  object
2   asins                               34658 non-null  object
3   brand                               34660 non-null  object
4   categories                           34660 non-null  object
5   keys                                 34660 non-null  object
6   manufacturer                         34660 non-null  object
7   reviews.date                        34621 non-null  object
8   reviews.dateAdded                   24039 non-null  object
9   reviews.dateSeen                    34660 non-null  object
10  reviews.didPurchase                 1 non-null      object
11  reviews.doRecommend                 34066 non-null  object
12  reviews.id                           1 non-null      float64
13  reviews.numHelpful                  34131 non-null  float64
14  reviews.rating                      34627 non-null  float64
15  reviews.sourceURLs                  34660 non-null  object
16  reviews.text                        34659 non-null  object
17  reviews.title                       34654 non-null  object
18  reviews.userCity                    0 non-null      float64
19  reviews.userProvince                0 non-null      float64
20  reviews.username                    34653 non-null  object
dtypes: float64(5), object(16)
memory usage: 5.6+ MB
```

In [8]: df.dtypes

Out[8]:

id	object
name	object
asins	object
brand	object
categories	object
keys	object
manufacturer	object
reviews.date	object
reviews.dateAdded	object
reviews.dateSeen	object
reviews.didPurchase	object
reviews.doRecommend	object
reviews.id	float64
reviews.numHelpful	float64
reviews.rating	float64
reviews.sourceURLs	object
reviews.text	object
reviews.title	object
reviews.userCity	float64
reviews.userProvince	float64
reviews.username	object

In [9]: df.tail()

Out[9]:

	id	name	asins	brand	categories	keys	manufacturer
34655	AVpfiBlyLJeJML43-4Tp	NaN	B006GWO5WK	Amazon	Computers/Tablets & Networking, Tablet & eBook ...	newamazonkindlefirehd9wpowerfastadaptercharger...	Amaz Digi Services, I
34656	AVpfiBlyLJeJML43-4Tp	NaN	B006GWO5WK	Amazon	Computers/Tablets & Networking, Tablet & eBook ...	newamazonkindlefirehd9wpowerfastadaptercharger...	Amaz Digi Services, I
34657	AVpfiBlyLJeJML43-4Tp	NaN	B006GWO5WK	Amazon	Computers/Tablets & Networking, Tablet & eBook ...	newamazonkindlefirehd9wpowerfastadaptercharger...	Amaz Digi Services, I
34658	AVpfiBlyLJeJML43-4Tp	NaN	B006GWO5WK	Amazon	Computers/Tablets & Networking, Tablet & eBook ...	newamazonkindlefirehd9wpowerfastadaptercharger...	Amaz Digi Services, I
34659	AVpfiBlyLJeJML43-4Tp	NaN	B006GWO5WK	Amazon	Computers/Tablets & Networking, Tablet & eBook ...	newamazonkindlefirehd9wpowerfastadaptercharger...	Amaz Digi Services, I

5 rows × 21 columns

In [10]:

df.describe()

Out[10]:

	reviews.id	reviews.numHelpful	reviews.rating	reviews.userCity	reviews.userProvince
count	1.0	34131.000000	34627.000000	0.0	0.0
mean	111372787.0	0.630248	4.584573	NaN	NaN
std	NaN	13.215775	0.735653	NaN	NaN
min	111372787.0	0.000000	1.000000	NaN	NaN
25%	111372787.0	0.000000	4.000000	NaN	NaN
50%	111372787.0	0.000000	5.000000	NaN	NaN
75%	111372787.0	0.000000	5.000000	NaN	NaN
max	111372787.0	814.000000	5.000000	NaN	NaN

In [11]:

columns\_to\_remove = ['reviews.userCity', 'reviews.userProvince', 'reviews.id']  
df = df.drop(columns=columns\_to\_remove, axis=1)

In [12]:

df

Out[12]:

		id	name	asins	brand	categories	keys	man
0	AVqklhwDv8e3D1O-lebb	All-New Fire HD 8 Tablet, 8 HD Display, Wi-Fi,...	B01AHB9CN2	Amazon	Electronics,iPad & Tablets,All Tablets,Fire Ta...	841667104676,amazon/53004484,amazon/b01ahb9cn2...		
1	AVqklhwDv8e3D1O-lebb	All-New Fire HD 8 Tablet, 8 HD Display, Wi-Fi,...	B01AHB9CN2	Amazon	Electronics,iPad & Tablets,All Tablets,Fire Ta...	841667104676,amazon/53004484,amazon/b01ahb9cn2...		
2	AVqklhwDv8e3D1O-lebb	All-New Fire HD 8 Tablet, 8 HD Display, Wi-Fi,...	B01AHB9CN2	Amazon	Electronics,iPad & Tablets,All Tablets,Fire Ta...	841667104676,amazon/53004484,amazon/b01ahb9cn2...		
3	AVqklhwDv8e3D1O-lebb	All-New Fire HD 8 Tablet, 8 HD Display, Wi-Fi,...	B01AHB9CN2	Amazon	Electronics,iPad & Tablets,All Tablets,Fire Ta...	841667104676,amazon/53004484,amazon/b01ahb9cn2...		
4	AVqklhwDv8e3D1O-lebb	All-New Fire HD 8 Tablet, 8 HD Display, Wi-Fi,...	B01AHB9CN2	Amazon	Electronics,iPad & Tablets,All Tablets,Fire Ta...	841667104676,amazon/53004484,amazon/b01ahb9cn2...		
...	...	...	...	...	...	...	...	...
34655	AVpfiBlyLJeJML43-4Tp	NaN	B006GWO5WK	Amazon	Computers/Tablets & Networking,Tablet & eBook ...	newamazonkindlefirehd9wpowerfastadaptercharger...		Se
34656	AVpfiBlyLJeJML43-4Tp	NaN	B006GWO5WK	Amazon	Computers/Tablets & Networking,Tablet & eBook ...	newamazonkindlefirehd9wpowerfastadaptercharger...		Se
34657	AVpfiBlyLJeJML43-4Tp	NaN	B006GWO5WK	Amazon	Computers/Tablets & Networking,Tablet & eBook ...	newamazonkindlefirehd9wpowerfastadaptercharger...		Se
34658	AVpfiBlyLJeJML43-4Tp	NaN	B006GWO5WK	Amazon	Computers/Tablets & Networking,Tablet & eBook ...	newamazonkindlefirehd9wpowerfastadaptercharger...		Se
34659	AVpfiBlyLJeJML43-4Tp	NaN	B006GWO5WK	Amazon	Computers/Tablets & Networking,Tablet & eBook ...	newamazonkindlefirehd9wpowerfastadaptercharger...		Se
34660 rows × 18 columns								
<div><div>◀</div><div></div><div>▶</div></div>								

## Dealing with missing values

```
In [13]: np.sum(df.isnull().any(axis=1))
```

Out[13]: 34659

```
In [14]: print('Count of columns in the data is: ', len(df.columns))
print('Count of rows in the data is: ', len(df))
```

Count of columns in the data is: 18  
Count of rows in the data is: 34660

```
In [15]: df['reviews.rating'].unique()
```

Out[15]: array([ 5., 4., 2., 1., 3., nan])

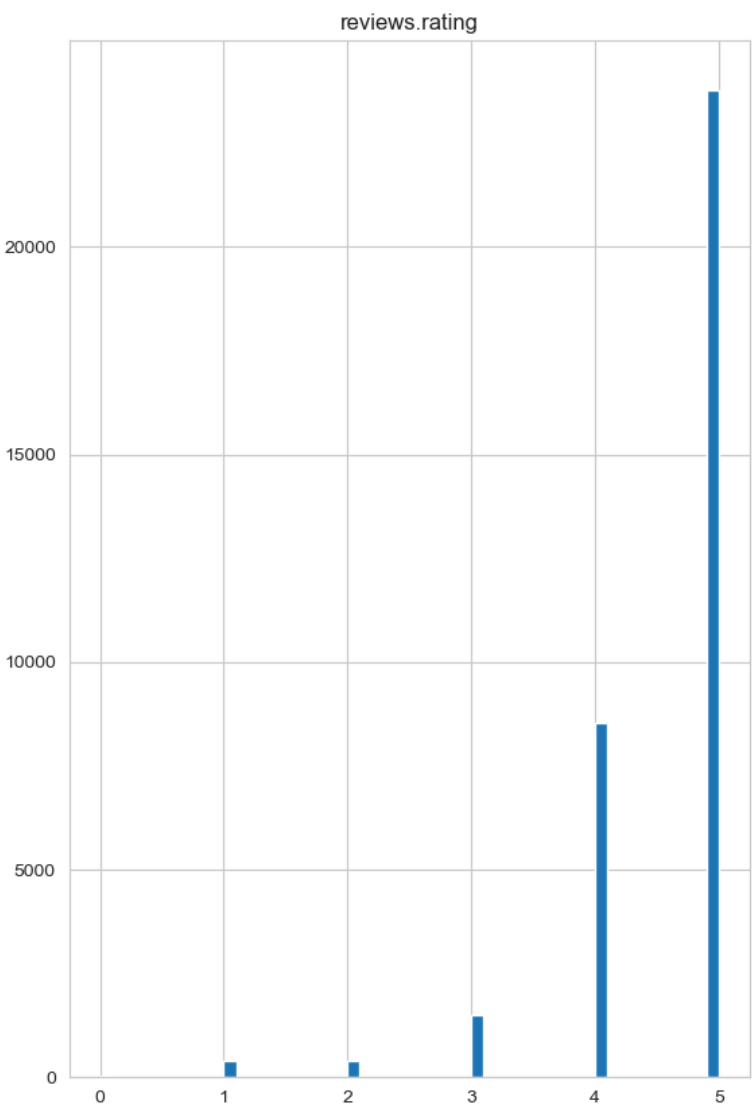

```
In [16]: df['reviews.rating'].nunique()
```

Out[16]: 5

```
In [17]: df['reviews.rating'].fillna(0, inplace=True)
print(df['reviews.rating'].unique())
```

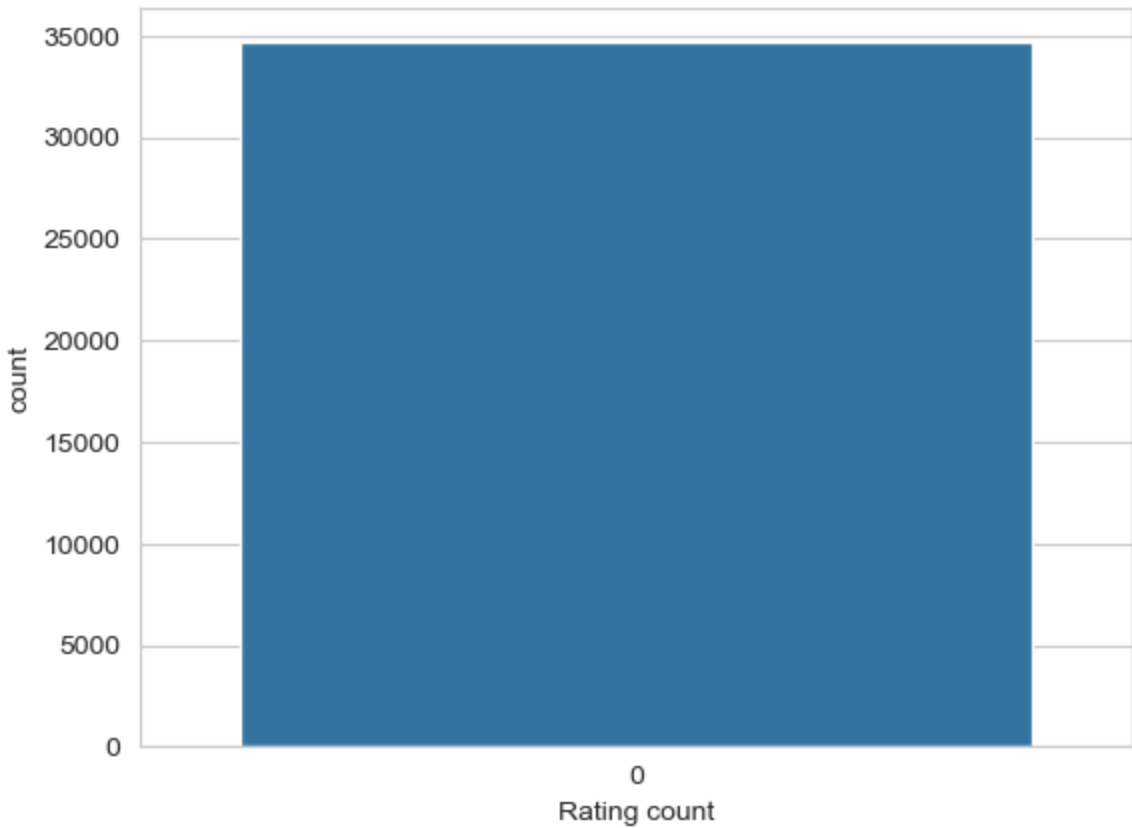
[5. 4. 2. 1. 3. 0.]

```
In [18]: df.hist(bins=50, figsize=(15,10))
plt.show()
```



```
In [19]: sns.countplot(df['reviews.rating'])
plt.xlabel('Rating count')
```

Out[19]: Text(0.5, 0, 'Rating count')



## Pre-processing Data

### using stop words

```
In [20]: from nltk.corpus import stopwords
df['reviews.text']=df['reviews.text'].str.lower()
# Download the stopwords corpus
nltk.download('stopwords')
```

[nltk\_data] Downloading package stopwords to  
[nltk\_data] C:\Users\User\AppData\Roaming\nltk\_data...  
[nltk\_data] Package stopwords is already up-to-date!

Out[20]: True

```
In [21]: stopwords_list = stopwords.words('english')
```

```
In [22]: from nltk.corpus import stopwords
", ".join(stopwords.words('english'))
```

Out[22]: "i, me, my, myself, we, our, ours, ourselves, you, you're, you've, you'll, you'd, your, yours, yourself, yo  
urselves, he, him, his, himself, she, she's, her, hers, herself, it, it's, its, itself, they, them, their,  
theirs, themselves, what, which, who, whom, this, that, that'll, these, those, am, is, are, was, were, be,  
been, being, have, has, had, having, do, does, did, doing, a, an, the, and, but, if, or, because, as, unti  
l, while, of, at, by, for, with, about, against, between, into, through, during, before, after, above, belo  
w, to, from, up, down, in, out, on, off, over, under, again, further, then, once, here, there, when, where,  
why, how, all, any, both, each, few, more, most, other, some, such, no, nor, not, only, own, same, so, tha  
n, too, very, s, t, can, will, just, don, don't, should, should've, now, d, ll, m, o, re, ve, y, ain, aren,  
aren't, couldn, couldn't, didn, didn't, doesn, doesn't, hadn, hadn't, hasn, hasn't, haven, haven't, isn, is  
n't, ma, mightn, mightn't, mustn, mustn't, needn, needn't, shan, shan't, shouldn, shouldn't, wasn, wasn't,  
weren, weren't, won, won't, wouldn, wouldn't"

```
In [23]: STOPWORDS = set(stopwords.words('english'))
def cleaning_stopwords(text):
    return " ".join([word for word in str(text).split() if word not in STOPWORDS])
df['reviews.text'] = df['reviews.text'].apply(lambda x: " ".join([word for word in str(x).split() if word.lo
df['reviews.text'].head()
```

Out[23]: 0 product far disappointed. children love use li...  
1 great beginner experienced person. bought gift...  
2 inexpensive tablet use learn on, step nabi. th...  
3 i've fire hd 8 two weeks love it. tablet great...  
4 bought grand daughter comes visit. set user, e...  
Name: reviews.text, dtype: object

```
In [24]: english_punctuations = string.punctuation
punctuations_list = english_punctuations
def cleaning_punctuations(text):
    translator = str.maketrans('', '', punctuations_list)
    return text.translate(translator)
```

```
In [25]: df['reviews.text'] = df['reviews.text'].apply(lambda x: cleaning_punctuations(x))
df['reviews.text'].tail()
```

```
Out[25]: 34655    appreciably faster 18 higher amp charger used ...
34656    amazon include charger kindle fact theyre char...
34657    love kindle fire really disappointed kindle po...
34658    surprised find come type charging cords purcha...
34659    spite fact nothing good things say amazon anth...
Name: reviews.text, dtype: object
```

## Using Stemming

```
In [26]: st = nltk.PorterStemmer()

def stemming_on_reviews_text(text):
    return [st.stem(word) for word in text]

df['reviews.text'] = df['reviews.text'].apply(lambda x: stemming_on_reviews_text(x))
```

```
In [27]: df['reviews.text'].head()
```

```
Out[27]: 0    [p, r, o, d, u, c, t, , f, a, r, , d, i, s, ...
1    [g, r, e, a, t, , b, e, g, i, n, n, e, r, , ...
2    [i, n, e, x, p, e, n, s, i, v, e, , t, a, b, ...
3    [i, v, e, , f, i, r, e, , h, d, , 8, , t, ...
4    [b, o, u, g, h, t, , g, r, a, n, d, , d, a, ...
Name: reviews.text, dtype: object
```

## Using Lemmatization

```
In [28]: lm = nltk.WordNetLemmatizer()
def lemmatizer_on_reviews_text(text):
    return[lm.lemmatize(word) for word in df]

df['reviews.text'] = df['reviews.text'].apply(lambda x: lemmatizer_on_reviews_text(x))
```

```
In [29]: df['reviews.text'].head()
```

```
Out[29]: 0    [id, name, asin, brand, category, key, manufac...
1    [id, name, asin, brand, category, key, manufac...
2    [id, name, asin, brand, category, key, manufac...
3    [id, name, asin, brand, category, key, manufac...
4    [id, name, asin, brand, category, key, manufac...
Name: reviews.text, dtype: object
```

In [30]:

df

Out[30]:

		id	name	asins	brand	categories	keys	man
0	AVqkIhwDv8e3D1O-lebb	All-New Fire HD 8 Tablet, 8 HD Display, Wi-Fi,...	B01AHB9CN2	Amazon	Electronics,iPad & Tablets,All Tablets,Fire Ta...	841667104676,amazon/53004484,amazon/b01ahb9cn2...		
1	AVqkIhwDv8e3D1O-lebb	All-New Fire HD 8 Tablet, 8 HD Display, Wi-Fi,...	B01AHB9CN2	Amazon	Electronics,iPad & Tablets,All Tablets,Fire Ta...	841667104676,amazon/53004484,amazon/b01ahb9cn2...		
2	AVqkIhwDv8e3D1O-lebb	All-New Fire HD 8 Tablet, 8 HD Display, Wi-Fi,...	B01AHB9CN2	Amazon	Electronics,iPad & Tablets,All Tablets,Fire Ta...	841667104676,amazon/53004484,amazon/b01ahb9cn2...		
3	AVqkIhwDv8e3D1O-lebb	All-New Fire HD 8 Tablet, 8 HD Display, Wi-Fi,...	B01AHB9CN2	Amazon	Electronics,iPad & Tablets,All Tablets,Fire Ta...	841667104676,amazon/53004484,amazon/b01ahb9cn2...		
4	AVqkIhwDv8e3D1O-lebb	All-New Fire HD 8 Tablet, 8 HD Display, Wi-Fi,...	B01AHB9CN2	Amazon	Electronics,iPad & Tablets,All Tablets,Fire Ta...	841667104676,amazon/53004484,amazon/b01ahb9cn2...		
...	...	...	...	...	...	...		
34655	AVpfiBlyLJeJML43-4Tp	NaN	B006GWO5WK	Amazon	Computers/Tablets & Networking,Tablet & eBook ...	newamazonkindlefirehd9wpowerfastadaptercharger...		Se
34656	AVpfiBlyLJeJML43-4Tp	NaN	B006GWO5WK	Amazon	Computers/Tablets & Networking,Tablet & eBook ...	newamazonkindlefirehd9wpowerfastadaptercharger...		Se
34657	AVpfiBlyLJeJML43-4Tp	NaN	B006GWO5WK	Amazon	Computers/Tablets & Networking,Tablet & eBook ...	newamazonkindlefirehd9wpowerfastadaptercharger...		Se
34658	AVpfiBlyLJeJML43-4Tp	NaN	B006GWO5WK	Amazon	Computers/Tablets & Networking,Tablet & eBook ...	newamazonkindlefirehd9wpowerfastadaptercharger...		Se
34659	AVpfiBlyLJeJML43-4Tp	NaN	B006GWO5WK	Amazon	Computers/Tablets & Networking,Tablet & eBook ...	newamazonkindlefirehd9wpowerfastadaptercharger...		Se

34660 rows × 18 columns

In [31]:

x= df['reviews.text']  
x.head()

Out[31]:

0 [id, name, asin, brand, category, key, manufac...  
1 [id, name, asin, brand, category, key, manufac...  
2 [id, name, asin, brand, category, key, manufac...  
3 [id, name, asin, brand, category, key, manufac...  
4 [id, name, asin, brand, category, key, manufac...  
Name: reviews.text, dtype: object

In [32]:

y= df['reviews.rating']  
y.tail()

Out[32]:

34655 3.0  
34656 1.0  
34657 1.0  
34658 1.0  
34659 1.0  
Name: reviews.rating, dtype: float64



```
In [33]: X = df['reviews.text'].astype(str)
```

## splitting Data traning= 0.7, testing 0.3

```
In [34]: # splitting Data for Training and Testing in two parts
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=21)
```

```
In [35]: y_train
```

```
Out[35]: 30618    5.0
31577    4.0
9981     5.0
32239    5.0
15223    5.0
...
16432    4.0
8964     5.0
5944     4.0
5327     1.0
15305    5.0
Name: reviews.rating, Length: 24262, dtype: float64
```

## Uni-gram for results using models

```
In [36]: #uni-gram
from sklearn.feature_extraction.text import TfidfVectorizer
vectorizer = TfidfVectorizer(ngram_range=(1,1))

# Training data
X_train = vectorizer.fit_transform(X_train)

# Testing data
X_test = vectorizer.transform(X_test)
```

## Making prediction on the test set

```
In [37]: # uni-gram
from sklearn.metrics import accuracy_score
from sklearn.ensemble import RandomForestClassifier
print("Random Forest Result")
rfc = RandomForestClassifier(n_estimators=100, random_state=52)
pred = rfc.fit(X_train, y_train).predict(X_test)
print(accuracy_score(y_test,pred))

from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score,classification_report
dt = DecisionTreeClassifier(random_state=50)
print("Decision Tree Result")
DecisionTree=dt.fit(X_train, y_train).predict(X_test)
print(accuracy_score(y_test,DecisionTree))

from sklearn.svm import SVC
print("Support Vector Machine Result")
svm = SVC(kernel='linear', C=2.0, random_state=52)
svm.fit(X_train,y_train)
y_pred=svm.predict(X_test)
print(accuracy_score(y_test,y_pred))

from sklearn.linear_model import LogisticRegression
lr=LogisticRegression()
print("Logistic Regression Result")
logisticRegresion=lr.fit(X_train, y_train).predict(X_test)
print(accuracy_score(y_test,logisticRegresion))
```

```
Random Forest Result
0.6906135795345258
Decision Tree Result
0.6906135795345258
Support Vector Machine Result
0.6906135795345258
Logistic Regression Result
0.6906135795345258
```

## Compute Classification report



```
In [38]: #uni-gram
print("Random Forest")
print(classification_report(y_test,pred))

print("Decision Tree")
print(classification_report(y_test,DecisionTree))

print("Support Vector Machine")
print(classification_report(y_test,y_pred))

print("Logistic Regression")
print(classification_report(y_test,logisticRegresion))
```

Random Forest				
	precision	recall	f1-score	support
0.0	0.00	0.00	0.00	13
1.0	0.00	0.00	0.00	122
2.0	0.00	0.00	0.00	115
3.0	0.00	0.00	0.00	442
4.0	0.00	0.00	0.00	2525
5.0	0.69	1.00	0.82	7181

accuracy			0.69	10398
macro avg	0.12	0.17	0.14	10398
weighted avg	0.48	0.69	0.56	10398

Decision Tree				
	precision	recall	f1-score	support
0.0	0.00	0.00	0.00	13
1.0	0.00	0.00	0.00	122
2.0	0.00	0.00	0.00	115
3.0	0.00	0.00	0.00	442
4.0	0.00	0.00	0.00	2525
5.0	0.69	1.00	0.82	7181

accuracy			0.69	10398
macro avg	0.12	0.17	0.14	10398
weighted avg	0.48	0.69	0.56	10398

Support Vector Machine				
	precision	recall	f1-score	support
0.0	0.00	0.00	0.00	13
1.0	0.00	0.00	0.00	122
2.0	0.00	0.00	0.00	115
3.0	0.00	0.00	0.00	442
4.0	0.00	0.00	0.00	2525
5.0	0.69	1.00	0.82	7181

accuracy			0.69	10398
macro avg	0.12	0.17	0.14	10398
weighted avg	0.48	0.69	0.56	10398

Logistic Regression				
	precision	recall	f1-score	support
0.0	0.00	0.00	0.00	13
1.0	0.00	0.00	0.00	122
2.0	0.00	0.00	0.00	115
3.0	0.00	0.00	0.00	442
4.0	0.00	0.00	0.00	2525
5.0	0.69	1.00	0.82	7181

accuracy			0.69	10398
macro avg	0.12	0.17	0.14	10398
weighted avg	0.48	0.69	0.56	10398

```
In [39]: X = df['reviews.text'].astype(str)
```

```
In [40]: # splitting for Training-Testing
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=21)
```

## bi-gram for results using models

```
In [41]: #bi-gram
from sklearn.feature_extraction.text import TfidfVectorizer
vectorizer = TfidfVectorizer(ngram_range=(2,2))

# Training Data
X_train = vectorizer.fit_transform(X_train)

# Testing Data
X_test = vectorizer.transform(X_test)
```

## Making prediction on the test set

```
In [42]: # bi-gram
from sklearn.metrics import accuracy_score
from sklearn.ensemble import RandomForestClassifier
print("Random Forest Result")
rfc = RandomForestClassifier(n_estimators=100, random_state=52)
pred = rfc.fit(X_train, y_train).predict(X_test)
print(accuracy_score(y_test,pred))

from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score,classification_report
dt = DecisionTreeClassifier(random_state=50)
print("Decision Tree Result")
DecisionTree=dt.fit(X_train, y_train).predict(X_test)
print(accuracy_score(y_test,DecisionTree))

from sklearn.svm import SVC
print("Support Vector Machine Result")
svm = SVC(kernel='linear', C=2.0, random_state=52, gamma=0.001)
svm.fit(X_train, y_train)
y_pred = svm.predict(X_test)
print(accuracy_score(y_test, y_pred))

from sklearn.linear_model import LogisticRegression
lr=LogisticRegression()
print("Logistic Regression Result")
logisticRegresion=lr.fit(X_train, y_train).predict(X_test)
print(accuracy_score(y_test,logisticRegresion))
```

```
Random Forest Result
0.6906135795345258
Decision Tree Result
0.6906135795345258
Support Vector Machine Result
0.6906135795345258
Logistic Regression Result
0.6906135795345258
```

## Compute Classification report

```
In [43]: #bi-gram
print("Random Forest")
print(classification_report(y_test,pred))

print("Decision Tree")
print(classification_report(y_test,DecisionTree))

print("Support Vector Machine")
print(classification_report(y_test,y_pred))

print("Logistic Regression")
print(classification_report(y_test,logisticRegresion))
```

Random Forest					
	precision	recall	f1-score	support	
0.0	0.00	0.00	0.00	13	
1.0	0.00	0.00	0.00	122	
2.0	0.00	0.00	0.00	115	
3.0	0.00	0.00	0.00	442	
4.0	0.00	0.00	0.00	2525	
5.0	0.69	1.00	0.82	7181	

accuracy			0.69	10398	
macro avg	0.12	0.17	0.14	10398	
weighted avg	0.48	0.69	0.56	10398	

Decision Tree					
	precision	recall	f1-score	support	
0.0	0.00	0.00	0.00	13	
1.0	0.00	0.00	0.00	122	
2.0	0.00	0.00	0.00	115	
3.0	0.00	0.00	0.00	442	
4.0	0.00	0.00	0.00	2525	
5.0	0.69	1.00	0.82	7181	

accuracy			0.69	10398	
macro avg	0.12	0.17	0.14	10398	
weighted avg	0.48	0.69	0.56	10398	

Support Vector Machine					
	precision	recall	f1-score	support	
0.0	0.00	0.00	0.00	13	
1.0	0.00	0.00	0.00	122	
2.0	0.00	0.00	0.00	115	
3.0	0.00	0.00	0.00	442	
4.0	0.00	0.00	0.00	2525	
5.0	0.69	1.00	0.82	7181	

accuracy			0.69	10398	
macro avg	0.12	0.17	0.14	10398	
weighted avg	0.48	0.69	0.56	10398	

Logistic Regression					
	precision	recall	f1-score	support	
0.0	0.00	0.00	0.00	13	
1.0	0.00	0.00	0.00	122	
2.0	0.00	0.00	0.00	115	
3.0	0.00	0.00	0.00	442	
4.0	0.00	0.00	0.00	2525	
5.0	0.69	1.00	0.82	7181	

accuracy			0.69	10398	
macro avg	0.12	0.17	0.14	10398	
weighted avg	0.48	0.69	0.56	10398	

```
In [44]: X = df['reviews.text'].astype(str)
```

```
In [45]: # splitting for Training-Testing
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=21)
```

## Tri-gram for results using models

```
In [46]: #Tri-gram
from sklearn.feature_extraction.text import TfidfVectorizer
vectorizer = TfidfVectorizer(ngram_range=(3,3))

# Training Data
X_train = vectorizer.fit_transform(X_train)

# Testing Data
X_test = vectorizer.transform(X_test)
```

## Making prediction on the test set

```
In [47]: # tri-gram
from sklearn.metrics import accuracy_score
from sklearn.ensemble import RandomForestClassifier
print("Random Forest Result")
rfc = RandomForestClassifier(n_estimators=100, random_state=52)
pred = rfc.fit(X_train, y_train).predict(X_test)
print(accuracy_score(y_test,pred))

from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score,classification_report
dt = DecisionTreeClassifier(random_state=50)
print("Decision Tree Result")
DecisionTree=dt.fit(X_train, y_train).predict(X_test)
print(accuracy_score(y_test,DecisionTree))

from sklearn.svm import SVC
print("Support Vector Machine Result")
svm = SVC(kernel='linear', C=2.0, random_state=52, gamma=0.001)
svm.fit(X_train, y_train)
y_pred = svm.predict(X_test)
print(accuracy_score(y_test, y_pred))

from sklearn.linear_model import LogisticRegression
lr=LogisticRegression()
print("Logistic Regression Result")
logisticRegression=lr.fit(X_train, y_train).predict(X_test)
print(accuracy_score(y_test,logisticRegresion))
```

```
Random Forest Result
0.6906135795345258
Decision Tree Result
0.6906135795345258
Support Vector Machine Result
0.6906135795345258
Logistic Regression Result
0.6906135795345258
```

## Compute Classification report

```
In [48]: #tri-gram
print("Random Forest")
print(classification_report(y_test,pred))

print("Decision Tree")
print(classification_report(y_test,DecisionTree))

print("Support Vector Machine")
print(classification_report(y_test,y_pred))

print("Logistic Regression")
print(classification_report(y_test,logisticRegresion))
```

Random Forest				
	precision	recall	f1-score	support
0.0	0.00	0.00	0.00	13
1.0	0.00	0.00	0.00	122
2.0	0.00	0.00	0.00	115
3.0	0.00	0.00	0.00	442
4.0	0.00	0.00	0.00	2525
5.0	0.69	1.00	0.82	7181

accuracy			0.69	10398
macro avg	0.12	0.17	0.14	10398
weighted avg	0.48	0.69	0.56	10398

Decision Tree				
	precision	recall	f1-score	support
0.0	0.00	0.00	0.00	13
1.0	0.00	0.00	0.00	122
2.0	0.00	0.00	0.00	115
3.0	0.00	0.00	0.00	442
4.0	0.00	0.00	0.00	2525
5.0	0.69	1.00	0.82	7181

accuracy			0.69	10398
macro avg	0.12	0.17	0.14	10398
weighted avg	0.48	0.69	0.56	10398

Support Vector Machine				
	precision	recall	f1-score	support
0.0	0.00	0.00	0.00	13
1.0	0.00	0.00	0.00	122
2.0	0.00	0.00	0.00	115
3.0	0.00	0.00	0.00	442
4.0	0.00	0.00	0.00	2525
5.0	0.69	1.00	0.82	7181

accuracy			0.69	10398
macro avg	0.12	0.17	0.14	10398
weighted avg	0.48	0.69	0.56	10398

Logistic Regression				
	precision	recall	f1-score	support
0.0	0.00	0.00	0.00	13
1.0	0.00	0.00	0.00	122
2.0	0.00	0.00	0.00	115
3.0	0.00	0.00	0.00	442
4.0	0.00	0.00	0.00	2525
5.0	0.69	1.00	0.82	7181

accuracy			0.69	10398
macro avg	0.12	0.17	0.14	10398
weighted avg	0.48	0.69	0.56	10398

```
In [49]: X = df['reviews.text'].astype(str)
```

```
In [50]: # splitting for Training-Testing
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=21)
```

## n-gram for results using models

```
In [51]: #n-gram
from sklearn.feature_extraction.text import TfidfVectorizer
vectorizer = TfidfVectorizer(ngram_range=(1,3))

# Training Data
X_train = vectorizer.fit_transform(X_train)

# Testing Data
X_test = vectorizer.transform(X_test)
```

## Making prediction on the test set

```
In [52]: # n-gram
from sklearn.metrics import accuracy_score
from sklearn.ensemble import RandomForestClassifier
print("Random Forest Result")
rfc = RandomForestClassifier(n_estimators=100, random_state=52)
pred = rfc.fit(X_train, y_train).predict(X_test)
print(accuracy_score(y_test,pred))

from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score,classification_report
dt = DecisionTreeClassifier(random_state=50)
print("Decision Tree Result")
DecisionTree=dt.fit(X_train, y_train).predict(X_test)
print(accuracy_score(y_test,DecisionTree))

from sklearn.svm import SVC
print("Support Vector Machine Result")
svm = SVC(kernel='linear', C=2.0, random_state=52)
svm.fit(X_train,y_train)
y_pred=svm.predict(X_test)
print(accuracy_score(y_test,y_pred))

from sklearn.linear_model import LogisticRegression
lr=LogisticRegression()
print("Logistic Regression Result")
logisticRegression=lr.fit(X_train, y_train).predict(X_test)
print(accuracy_score(y_test,logisticRegresion))
```

```
Random Forest Result
0.6906135795345258
Decision Tree Result
0.6906135795345258
Support Vector Machine Result
0.6906135795345258
Logistic Regression Result
0.6906135795345258
```

## Compute Classification report

```
In [53]: #n-gram
print("Decision Tree")
print(classification_report(y_test,DecisionTree))

print("Random Forest")
print(classification_report(y_test,pred))

print("Logistic Regression")
print(classification_report(y_test,logisticRegresion))

print("Support Vector Machine")
print(classification_report(y_test,y_pred))
```

Decision Tree					
	precision	recall	f1-score	support	
0.0	0.00	0.00	0.00	13	
1.0	0.00	0.00	0.00	122	
2.0	0.00	0.00	0.00	115	
3.0	0.00	0.00	0.00	442	
4.0	0.00	0.00	0.00	2525	
5.0	0.69	1.00	0.82	7181	
accuracy			0.69	10398	
macro avg	0.12	0.17	0.14	10398	
weighted avg	0.48	0.69	0.56	10398	
Random Forest					
	precision	recall	f1-score	support	
0.0	0.00	0.00	0.00	13	
1.0	0.00	0.00	0.00	122	
2.0	0.00	0.00	0.00	115	
3.0	0.00	0.00	0.00	442	
4.0	0.00	0.00	0.00	2525	
5.0	0.69	1.00	0.82	7181	
accuracy			0.69	10398	
macro avg	0.12	0.17	0.14	10398	
weighted avg	0.48	0.69	0.56	10398	
Logistic Regression					
	precision	recall	f1-score	support	
0.0	0.00	0.00	0.00	13	
1.0	0.00	0.00	0.00	122	
2.0	0.00	0.00	0.00	115	
3.0	0.00	0.00	0.00	442	
4.0	0.00	0.00	0.00	2525	
5.0	0.69	1.00	0.82	7181	
accuracy			0.69	10398	
macro avg	0.12	0.17	0.14	10398	
weighted avg	0.48	0.69	0.56	10398	
Support Vector Machine					
	precision	recall	f1-score	support	
0.0	0.00	0.00	0.00	13	
1.0	0.00	0.00	0.00	122	
2.0	0.00	0.00	0.00	115	
3.0	0.00	0.00	0.00	442	
4.0	0.00	0.00	0.00	2525	
5.0	0.69	1.00	0.82	7181	
accuracy			0.69	10398	
macro avg	0.12	0.17	0.14	10398	
weighted avg	0.48	0.69	0.56	10398	

```
In [54]: # Assuming df is your cleaned DataFrame
df.to_csv('amazon_data.csv', index=False)
```

In [ ]: